

Machine Learning Engineer Nanodegree

Unsupervised Learning

Project: Creating Customer Segments

Welcome to the third project of the Machine Learning Engineer Nanodegree! In this notebook, some template code has already been provided for you, and it will be your job to implement the additional functionality necessary to successfully complete this project. Sections that begin with **'Implementation'** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section and the specifics of the implementation are marked in the code block with a `'TODO'` statement. Please be sure to read the instructions carefully!

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question X'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

Getting Started

In this project, you will analyze a dataset containing data on various customers' annual spending amounts (reported in *monetary units*) of diverse product categories for internal structure. One goal of this project is to best describe the variation in the different types of customers that a wholesale distributor interacts with. Doing so would equip the distributor with insight into how to best structure their delivery service to meet the needs of each customer.

The dataset for this project can be found on the [UCI Machine Learning Repository \(https://archive.ics.uci.edu/ml/datasets/Wholesale+customers\)](https://archive.ics.uci.edu/ml/datasets/Wholesale+customers). For the purposes of this project, the features `'Channel'` and `'Region'` will be excluded in the analysis — with focus instead on the six product categories recorded for customers.

Run the code block below to load the wholesale customers dataset, along with a few of the necessary Python libraries required for this project. You will know the dataset loaded successfully if the size of the dataset is reported.

```
In [1]: 1 # Import libraries necessary for this project
2 import numpy as np
3 import pandas as pd
4 from IPython.display import display # Allows the use of display() for DataFrames
5
6 # Import supplementary visualizations code visuals.py
7 import visuals as vs
8
9 # Pretty display for notebooks
10 %matplotlib inline
11
12 # Load the wholesale customers dataset
13 try:
14     data = pd.read_csv("customers.csv")
15     data.drop(['Region', 'Channel'], axis = 1, inplace = True)
16     print "Wholesale customers dataset has {} samples with {} features each.".format(*data.shape)
17 except:
18     print "Dataset could not be loaded. Is the dataset missing?"
```

Wholesale customers dataset has 440 samples with 6 features each.

Data Exploration

In this section, you will begin exploring the data through visualizations and code to understand how each feature is related to the others. You will observe a statistical description of the dataset, consider the relevance of each feature, and select a few sample data points from the dataset which you will track through the course of this project.

Run the code block below to observe a statistical description of the dataset. Note that the dataset is composed of six important product categories: **'Fresh'**, **'Milk'**, **'Grocery'**, **'Frozen'**, **'Detergents_Paper'**, and **'Delicatessen'**. Consider what each category represents in terms of products you could purchase.

In [2]:

```
1 # Display a description of the dataset
2 display(data.describe())
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
count	440.000000	440.000000	440.000000	440.000000	440.000000	440.000000
mean	12000.297727	5796.265909	7951.277273	3071.931818	2881.493182	1524.870455
std	12647.328865	7380.377175	9503.162829	4854.673333	4767.854448	2820.105937
min	3.000000	55.000000	3.000000	25.000000	3.000000	3.000000
25%	3127.750000	1533.000000	2153.000000	742.250000	256.750000	408.250000
50%	8504.000000	3627.000000	4755.500000	1526.000000	816.500000	965.500000
75%	16933.750000	7190.250000	10655.750000	3554.250000	3922.000000	1820.250000
max	112151.000000	73498.000000	92780.000000	60869.000000	40827.000000	47943.000000

Implementation: Selecting Samples

To get a better understanding of the customers and how their data will transform through the analysis, it would be best to select a few sample data points and explore them in more detail. In the code block below, add **three** indices of your choice to the `indices` list which will represent the customers to track. It is suggested to try different sets of samples until you obtain customers that vary significantly from one another.

In [3]:

```
1 # TODO: Select three indices of your choice you wish to sample from the dataset
2 indices = [0, 1, 439]
3
4 # Create a DataFrame of the chosen samples
5 samples = pd.DataFrame(data.loc[indices], columns = data.keys()).reset_index(drop = True)
6 print "Chosen samples of wholesale customers dataset:"
7 display(samples)
```

Chosen samples of wholesale customers dataset:

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
0	12669	9656	7561	214	2674	1338
1	7057	9810	9568	1762	3293	1776
2	2787	1698	2510	65	477	52

In [4]:

```
1 # More insights on the chosen samples
2 display((samples - data.mean())/data.std())
3 display((samples - data.median())/data.std())
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
0	0.052873	0.522972	-0.041068	-0.588697	-0.043519	-0.066264
1	-0.390857	0.543839	0.170125	-0.269829	0.086309	0.089050
2	-0.728478	-0.555292	-0.572575	-0.619389	-0.504313	-0.522275

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
0	0.329319	0.816896	0.295218	-0.270255	0.389588	0.132087
1	-0.114412	0.837762	0.506410	0.048613	0.519416	0.287401
2	-0.452032	-0.261369	-0.236290	-0.300947	-0.071206	-0.323924

Question 1

Consider the total purchase cost of each product category and the statistical description of the dataset above for your sample customers.

- What kind of establishment (customer) could each of the three samples you've chosen represent?

Hint: Examples of establishments include places like markets, cafes, delis, wholesale retailers, among many others. Avoid using names for establishments, such as saying *"McDonalds"* when describing a sample customer as a restaurant. You can use the mean values for reference to compare your samples with. The mean values are as follows:

- Fresh: 12000.2977
- Milk: 5796.2
- Grocery: 3071.9
- Detergents_paper: 2881.4
- Delicatessen: 1524.8

Knowing this, how do your samples compare? Does that help in driving your insight into what kind of establishments they might be?

Answer: Refer to the tabble from previous section. It shows how many sigma far from the mean and median of each product. The first and second samples have similar spendings on almost all the products except 'Fresh' and 'Frozen'. Sample 0 consumes more 'Fresh' than sample 1, 0.45 sigma, while 0.32 sigma of 'Frozen'. Sample 2 has much less spending on all the products, with 0.5 to 0.7 sigma less than average. All the observations should help distinguish them from each other.

Maybe the first one is cafe, second one is delis and the third one is small restaurant.

Implementation: Feature Relevance

One interesting thought to consider is if one (or more) of the six product categories is actually relevant for understanding customer purchasing. That is to say, is it possible to determine whether customers purchasing some amount of one category of products will necessarily purchase some proportional amount of another category of products? We can make this determination quite easily by training a supervised regression learner on a subset of the data with one feature removed, and then score how well that model can predict the removed feature.

In the code block below, you will need to implement the following:

- Assign `new_data` a copy of the data by removing a feature of your choice using the `DataFrame.drop` function.
- Use `sklearn.cross_validation.train_test_split` to split the dataset into training and testing sets.
 - Use the removed feature as your target label. Set a `test_size` of 0.25 and set a `random_state`.
- Import a decision tree regressor, set a `random_state`, and fit the learner to the training data.
- Report the prediction score of the testing set using the regressor's `score` function.

```
In [5]: 1 # TODO: Make a copy of the DataFrame, using the 'drop' function to drop the given feature
2 new_data = data.drop('Detergents_Paper', axis=1)
3 new_target = data['Detergents_Paper']
4
5 # TODO: Split the data into training and testing sets(0.25) using the given feature as the target
6 # Set a random state.
7 from sklearn.model_selection import train_test_split
8 X_train, X_test, y_train, y_test = train_test_split(new_data, new_target, test_size=0.25, random_state=0)
9
10 # TODO: Create a decision tree regressor and fit it to the training set
11 from sklearn.tree import DecisionTreeRegressor
12 regressor = DecisionTreeRegressor(random_state=0)
13 regressor.fit(X_train, y_train)
14
15 # TODO: Report the score of the prediction using the testing set
16 # The score can also be obtained when we set random_state from 0 to 100 and take the average
17 score = regressor.score(X_test, y_test)
18 print(score)
```

0.7286551812541454

Question 2

- Which feature did you attempt to predict?
- What was the reported prediction score?
- Is this feature necessary for identifying customers' spending habits?


Hint: The coefficient of determination, R^2 , is scored between 0 and 1, with 1 being a perfect fit. A negative R^2 implies the model fails to fit the data. If you get a low score for a particular feature, that lends us to beleive that that feature point is hard to predict using the other features, thereby making it an important feature to consider when considering relevance.

Answer: I chose 'Detergents_Paper' as the feature to predict. The reported R2 score is 0.73, not a very good sign on the 'fitness', but also not bad. Actually it's the predicting feature that creates the biggest R2 score. Given by the score, I will treat it as an unnecessary feature.

Visualize Feature Distributions


To get a better understanding of the dataset, we can construct a scatter matrix of each of the six product features present in the data. If you found that the feature you attempted to predict above is relevant for identifying a specific customer, then the scatter matrix below may not show any correlation between that feature and the others. Conversely, if you believe that feature is not relevant for identifying a specific customer, the scatter matrix might show a correlation between that feature and another feature in the data. Run the code block below to produce a scatter matrix.

```
In [6]: 1 # Produce a scatter matrix for each pair of features in the data
2 pd.plotting.scatter_matrix(data, alpha = 0.3, figsize = (14,8), diagonal = 'kde');
```



```
In [7]: 1 import seaborn as sns
2 sns.heatmap(data.corr())
```

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1003d4e0>



Question 3

- Using the scatter matrix as a reference, discuss the distribution of the dataset, specifically talk about the normality, outliers, large number of data points near 0 among others. If you need to separate out some of the plots individually to further accentuate your point, you may do so as well.
- Are there any pairs of features which exhibit some degree of correlation?
- Does this confirm or deny your suspicions about the relevance of the feature you attempted to predict?
- How is the data for those features distributed?

Hint: Is the data normally distributed? Where do most of the data points lie? You can use `corr()` (<https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.corr.html>) to get the feature correlations and then visualize them using a [heatmap](http://seaborn.pydata.org/generated/seaborn.heatmap.html) (<http://seaborn.pydata.org/generated/seaborn.heatmap.html>) (the data that would be fed into the heatmap would be the correlation values, for eg: `data.corr()`) to gain further insight.

Answer: In the scatter matrix, majority of the data points are near 0. Also few data points are closed to the upper limit of the value range.

The pair of **'Grocery'** and **'Detergents_Paper'** has a relatively high degree of correlation. This can also be verified from the heatmap.

In my view, none of the features has data normally distributed. But due to few outliers, the distributions are skewed. Maybe some preprocessing, like taking a logarithmic could be a good choice.

Data Preprocessing

In this section, you will preprocess the data to create a better representation of customers by performing a scaling on the data and detecting (and optionally removing) outliers. Preprocessing data is often times a critical step in assuring that results you obtain from your analysis are significant and meaningful.

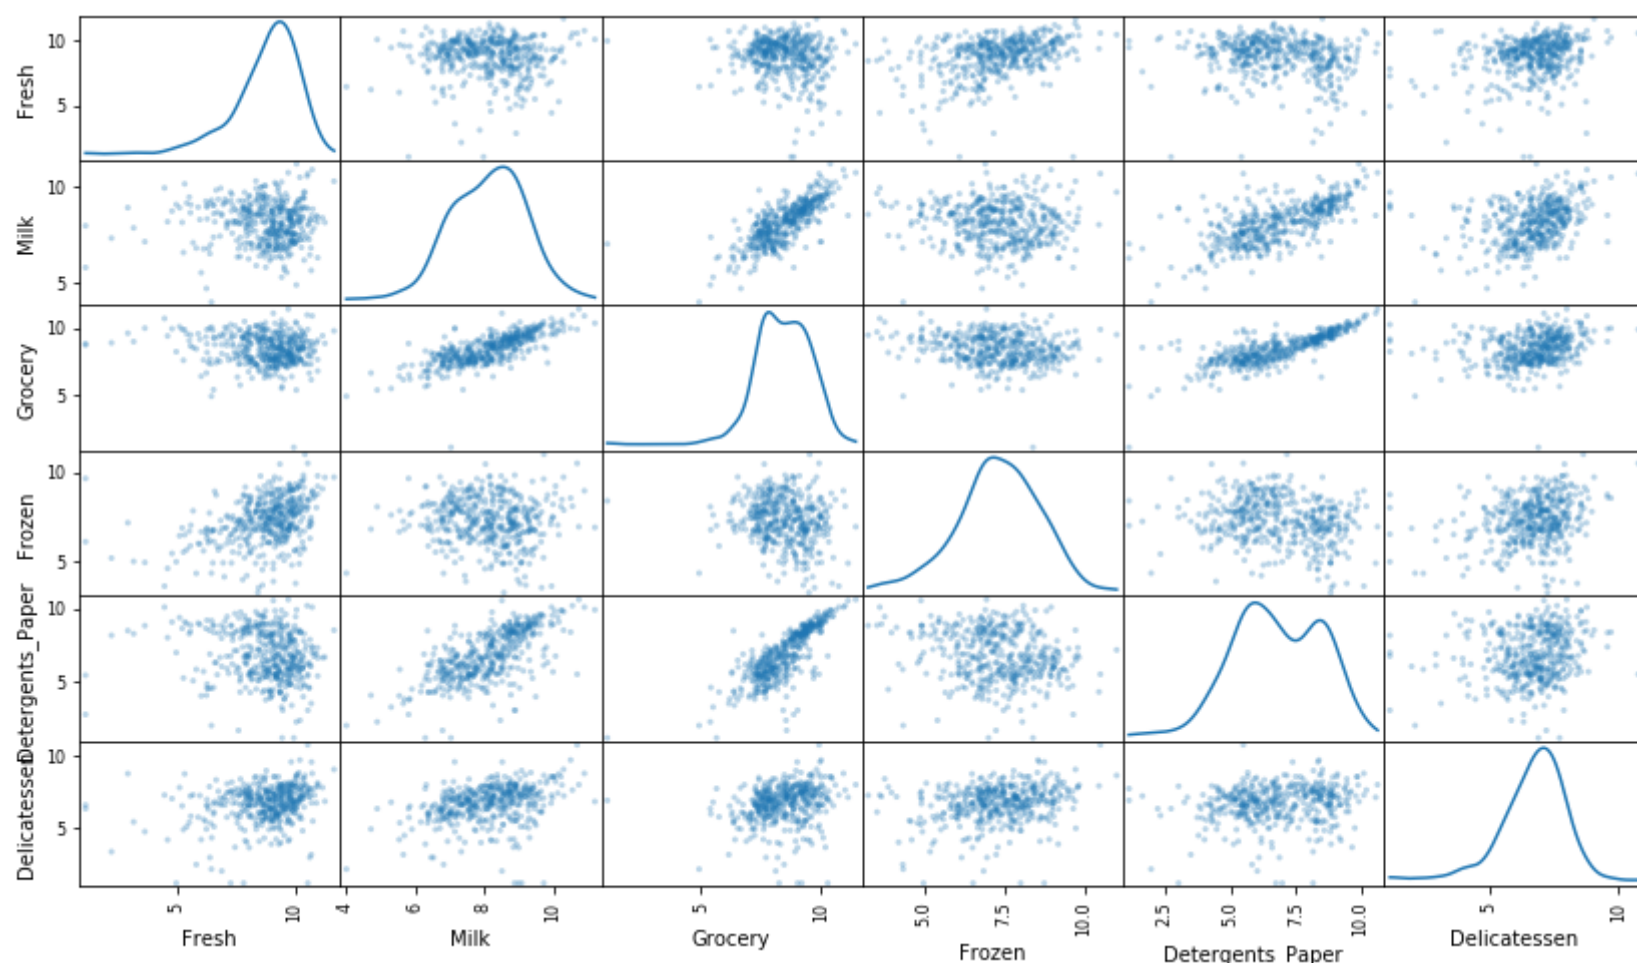
Implementation: Feature Scaling

If data is not normally distributed, especially if the mean and median vary significantly (indicating a large skew), it is most [often appropriate](http://econbrowser.com/archives/2014/02/use-of-logarithms-in-economics) (<http://econbrowser.com/archives/2014/02/use-of-logarithms-in-economics>) to apply a non-linear scaling — particularly for financial data. One way to achieve this scaling is by using a [Box-Cox test](http://scipy.github.io/devdocs/generated/scipy.stats.boxcox.html) (<http://scipy.github.io/devdocs/generated/scipy.stats.boxcox.html>), which calculates the best power transformation of the data that reduces skewness. A simpler approach which can work in most cases would be applying the natural logarithm.

In the code block below, you will need to implement the following:

- Assign a copy of the data to `log_data` after applying logarithmic scaling. Use the `np.log` function for this.
- Assign a copy of the sample data to `log_samples` after applying logarithmic scaling. Again, use `np.log`.

```
In [8]: 1 # TODO: Scale the data using the natural logarithm
2 log_data = np.log(data)
3
4 # TODO: Scale the sample data using the natural logarithm
5 log_samples = np.log(samples)
6
7 # Produce a scatter matrix for each pair of newly-transformed features
8 pd.plotting.scatter_matrix(log_data, alpha = 0.3, figsize = (14,8), diagonal = 'kde');
```



Observation

After applying a natural logarithm scaling to the data, the distribution of each feature should appear much more normal. For any pairs of features you may have identified earlier as being correlated, observe here whether that correlation is still present (and whether it is now stronger or weaker than before).

Run the code below to see how the sample data has changed after having the natural logarithm applied to it.

In [9]:

1

Display the log-transformed sample data

2

display(log_samples)

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
0	9.446913	9.175335	8.930759	5.365976	7.891331	7.198931
1	8.861775	9.191158	9.166179	7.474205	8.099554	7.482119
2	7.932721	7.437206	7.828038	4.174387	6.167516	3.951244

Implementation: Outlier Detection

Detecting outliers in the data is extremely important in the data preprocessing step of any analysis. The presence of outliers can often skew results which take into consideration these data points. There are many "rules of thumb" for what constitutes an outlier in a dataset. Here, we will use [Tukey's Method for identifying outliers \(http://datapigtechnologies.com/blog/index.php/highlighting-outliers-in-your-data-with-the-tukey-method/\)](http://datapigtechnologies.com/blog/index.php/highlighting-outliers-in-your-data-with-the-tukey-method/): An *outlier step* is calculated as 1.5 times the interquartile range (IQR). A data point with a feature that is beyond an outlier step outside of the IQR for that feature is considered abnormal.

In the code block below, you will need to implement the following:

- Assign the value of the 25th percentile for the given feature to `Q1` . Use `np.percentile` for this.
- Assign the value of the 75th percentile for the given feature to `Q3` . Again, use `np.percentile` .
- Assign the calculation of an outlier step for the given feature to `step` .
- Optionally remove data points from the dataset by adding indices to the `outliers` list.

NOTE: If you choose to remove any outliers, ensure that the sample data does not contain any of these points! Once you have performed this implementation, the dataset will be stored in the variable `good_data` .

In [10]:

```
1 # For each feature find the data points with extreme high or low values
2 temp_outliers = []
3 for feature in log_data.keys():
4
5     # TODO: Calculate Q1 (25th percentile of the data) for the given feature
6     Q1 = np.percentile(log_data[feature], 25)
7
8     # TODO: Calculate Q3 (75th percentile of the data) for the given feature
9     Q3 = np.percentile(log_data[feature], 75)
10
11     # TODO: Use the interquartile range to calculate an outlier step (1.5 times the interquartile range)
12     step = 1.5 * (Q3 - Q1)
13
14     # Display the outliers
15     print "Data points considered outliers for the feature '{}':".format(feature)
16     display(log_data[~((log_data[feature] >= Q1 - step) & (log_data[feature] <= Q3 + step))])
17
18     # Get the index of lower and upper outliers, the first element of the np.where()
19     lower_outliers_indices = np.where(log_data[feature] < Q1-step)[0]
20     upper_outliers_indices = np.where(log_data[feature] > Q3+step)[0]
21
22     temp_outliers.append(lower_outliers_indices)
23     temp_outliers.append(upper_outliers_indices)
24
25 # OPTIONAL: Select the indices for data points you wish to remove
26 outliers = []
27 for i in temp_outliers:
28     outliers = outliers + i.tolist()
29
30 import collections
31 cnt = collections.Counter(outliers)
32 num_dup = 0
33 for i in cnt.viewvalues():
34     if i > 1:
35         num_dup += 1
36 print("There are " + str(num_dup) + " duplicates.")
37
38 outliers = list(set(outliers))
39 print("There are " + str(len(outliers)) + " outliers.")
40
41 # Remove the outliers, if any were specified
42 good_data = log_data.drop(log_data.index[outliers]).reset_index(drop = True)
```

Data points considered outliers for the feature 'Fresh':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
65	4.442651	9.950323	10.732651	3.583519	10.095388	7.260523
66	2.197225	7.335634	8.911530	5.164786	8.151333	3.295837
81	5.389072	9.163249	9.575192	5.645447	8.964184	5.049856
95	1.098612	7.979339	8.740657	6.086775	5.407172	6.563856
96	3.135494	7.869402	9.001839	4.976734	8.262043	5.379897
128	4.941642	9.087834	8.248791	4.955827	6.967909	1.098612
171	5.298317	10.160530	9.894245	6.478510	9.079434	8.740337
193	5.192957	8.156223	9.917982	6.865891	8.633731	6.501290
218	2.890372	8.923191	9.629380	7.158514	8.475746	8.759669
304	5.081404	8.917311	10.117510	6.424869	9.374413	7.787382
305	5.493061	9.468001	9.088399	6.683361	8.271037	5.351858
338	1.098612	5.808142	8.856661	9.655090	2.708050	6.309918
353	4.762174	8.742574	9.961898	5.429346	9.069007	7.013016
355	5.247024	6.588926	7.606885	5.501258	5.214936	4.844187
357	3.610918	7.150701	10.011086	4.919981	8.816853	4.700480
412	4.574711	8.190077	9.425452	4.584967	7.996317	4.127134

Data points considered outliers for the feature 'Milk':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
86	10.039983	11.205013	10.377047	6.894670	9.906981	6.805723
98	6.220590	4.718499	6.656727	6.796824	4.025352	4.882802
154	6.432940	4.007333	4.919981	4.317488	1.945910	2.079442
356	10.029503	4.897840	5.384495	8.057377	2.197225	6.306275

Data points considered outliers for the feature 'Grocery':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
--	-------	------	---------	--------	------------------	--------------

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
75	9.923192	7.036148	1.098612	8.390949	1.098612	6.882437
154	6.432940	4.007333	4.919981	4.317488	1.945910	2.079442

Data points considered outliers for the feature 'Frozen':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
38	8.431853	9.663261	9.723703	3.496508	8.847360	6.070738
57	8.597297	9.203618	9.257892	3.637586	8.932213	7.156177
65	4.442651	9.950323	10.732651	3.583519	10.095388	7.260523
145	10.000569	9.034080	10.457143	3.737670	9.440738	8.396155
175	7.759187	8.967632	9.382106	3.951244	8.341887	7.436617
264	6.978214	9.177714	9.645041	4.110874	8.696176	7.142827
325	10.395650	9.728181	9.519735	11.016479	7.148346	8.632128
420	8.402007	8.569026	9.490015	3.218876	8.827321	7.239215
429	9.060331	7.467371	8.183118	3.850148	4.430817	7.824446
439	7.932721	7.437206	7.828038	4.174387	6.167516	3.951244

Data points considered outliers for the feature 'Detergents_Paper':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
75	9.923192	7.036148	1.098612	8.390949	1.098612	6.882437
161	9.428190	6.291569	5.645447	6.995766	1.098612	7.711101

Data points considered outliers for the feature 'Delicatessen':

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
66	2.197225	7.335634	8.911530	5.164786	8.151333	3.295837
109	7.248504	9.724899	10.274568	6.511745	6.728629	1.098612
128	4.941642	9.087834	8.248791	4.955827	6.967909	1.098612
137	8.034955	8.997147	9.021840	6.493754	6.580639	3.583519
142	10.519646	8.875147	9.018332	8.004700	2.995732	1.098612
154	6.432940	4.007333	4.919981	4.317488	1.945910	2.079442
183	10.514529	10.690808	9.911952	10.505999	5.476464	10.777768
184	5.789960	6.822197	8.457443	4.304065	5.811141	2.397895
187	7.798933	8.987447	9.192075	8.743372	8.148735	1.098612
203	6.368187	6.529419	7.703459	6.150603	6.860664	2.890372
233	6.871091	8.513988	8.106515	6.842683	6.013715	1.945910
285	10.602965	6.461468	8.188689	6.948897	6.077642	2.890372
289	10.663966	5.655992	6.154858	7.235619	3.465736	3.091042
343	7.431892	8.848509	10.177932	7.283448	9.646593	3.610918

There are 5 duplicates.
There are 42 outliers.

Question 4

- Are there any data points considered outliers for more than one feature based on the definition above?
- Should these data points be removed from the dataset?
- If any data points were added to the outliers list to be removed, explain why.

** Hint: ** If you have datapoints that are outliers in multiple categories think about why that may be and if they warrant removal. Also note how k-means is affected by outliers and whether or not this plays a factor in your analysis of whether or not to remove them.

Answer: Yes, there are 5 outliers exist in multiple features. In total, 42 outliers.

Yes, they should be removed from the dataset. In this case, the outliers may be due to wrong inputs, but also could be due to the customers are really big and important customers. When we make business decisions, VIPs may need to be treated seperately. Moreover, they will make the dataset skewed and impact the accuracy of the model for a general case.

Feature Transformation

In this section you will use principal component analysis (PCA) to draw conclusions about the underlying structure of the wholesale customer data. Since using PCA on a dataset calculates the dimensions which best maximize variance, we will find which compound combinations of features best describe customers.


Implementation: PCA

Now that the data has been scaled to a more normal distribution and has had any necessary outliers removed, we can now apply PCA to the `good_data` to discover which dimensions about the data best maximize the variance of features involved. In addition to finding these dimensions, PCA will also report the *explained variance ratio* of each dimension — how much variance within the data is explained by that dimension alone. Note that a component (dimension) from PCA can be considered a new "feature" of the space, however it is a composition of the original features present in the data.

In the code block below, you will need to implement the following:

- Import `sklearn.decomposition.PCA` and assign the results of fitting PCA in six dimensions with `good_data` to `pca`.
- Apply a PCA transformation of `log_samples` using `pca.transform`, and assign the results to `pca_samples`.

```
In [11]: 1 # TODO: Apply PCA by fitting the good data with the same number of dimensions as features
2 from sklearn.decomposition import PCA
3 pca = PCA(n_components=6)
4 pca.fit(good_data)
5
6 # TODO: Transform log_samples using the PCA fit above
7 pca_samples = pca.transform(log_samples)
8
9 # Generate PCA results plot
10 pca_results = vs.pca_results(good_data, pca)
11 print(np.cumsum(pca.explained_variance_ratio_))
```



Question 5

- How much variance in the data is explained* **in total** *by the first and second principal component?
- How much variance in the data is explained by the first four principal components?
- Using the visualization provided above, talk about each dimension and the cumulative variance explained by each, stressing upon which features are well represented by each dimension(both in terms of positive and negative variance explained). Discuss what the first four dimensions best represent in terms of customer spending.

Hint: A positive increase in a specific dimension corresponds with an *increase* of the *positive-weighted* features and a *decrease* of the *negative-weighted* features. The rate of increase or decrease is based on the individual feature weights.

Answer: 73% of variance in the data is explained the first two principal components. 93% of variance explained by the first four principal components.

In the first principal component, 'Milk', 'Grocery' and 'Detergents_Paper' contributed the most positive weights. The rest 3 products accounted smaller weights, 'Fresh' and 'Frozen' have negative weights in this scenario. This indicates that customers with large positive value of PC1 tends to buy more 'Milk', 'Grocery' and 'Detergents_Paper', they also tend to buy less 'Fresh' and 'Frozen', but the magnitude of 'less' is smaller. Because the 2 groups of products are inversely correlated. In another way to say, customers with large negative value of P1 tend to buy more 'Fresh' and 'Frozen', less the other productes. Almost half of the variance can be represented by the first pricipal component.

In the second principal component, only 'Detergents_Paper' is anti-correlated to the reset of products, but the manitude is very small. 'Fresh', 'Frozen' and 'Delicatessen' are high correlated, meaning if a customer buys more 'Fresh', he/she tends to buy more 'Frozen' and 'Delicatessen'. It's also true that he/she buys more 'Milk' and 'Grocery', but the spending increment may not be as much as 'Frozen' and 'Delicatessen'. Since 'Detergents_Paper' is anti-correlated to the rest of products, so this specific customer tends to buy slightly less 'Detergetns_Paper'.

The third principal component has 'Fresh' and 'Detergents_Paper' are highly correlated. 'Milk', 'Frozen' and 'Delicatessen' are highly correlated between each other, but anti-correlated from the earlier group of 'Fresh' and 'Detergents_Paper'. The magnitude of 'Grocery' is so small, I tend to ignore it in the analysis. So when a customer spends more in one product, he/she tends to spend more in the product from the same group which has high correlation between each other. But he/she spends less in the products from the group that anti-correlate to it.

The same analysis hold true for the fourth pricipal component. When a customer buys more 'Frozen', he/she tends to buy more 'Detergents_Paper', in the meantime, much less 'Delicatessen', a little less 'Fresh', 'Milk' and 'Grocery'.

Observation

Run the code below to see how the log-transformed sample data has changed after having a PCA transformation applied to it in six dimensions. Observe the numerical value for the first four dimensions of the sample points. Consider if this is consistent with your initial interpretation of the sample points.

```
In [12]: 1 # Display sample Log-data after having a PCA transformation applied
2 display(pd.DataFrame(np.round(pca_samples, 4), columns = pca_results.index.values))
```

	Dimension 1	Dimension 2	Dimension 3	Dimension 4	Dimension 5	Dimension 6
0	1.7793	-0.6604	-0.8247	1.6580	0.0903	-0.4808
1	1.8820	0.4617	0.2764	0.1055	0.0958	-0.2093
2	-0.9538	-4.0931	-1.6164	0.5626	0.5658	-0.0773

Implementation: Dimensionality Reduction

When using principal component analysis, one of the main goals is to reduce the dimensionality of the data — in effect, reducing the complexity of the problem. Dimensionality reduction comes at a cost: Fewer dimensions used implies less of the total variance in the data is being explained. Because of this, the *cumulative explained variance ratio* is extremely important for knowing how many dimensions are necessary for the problem. Additionally, if a signifiant amount of variance is explained by only two or three dimensions, the reduced data can be visualized afterwards.

In the code block below, you will need to implement the following:

- Assign the results of fitting PCA in two dimensions with `good_data` to `pca` .
- Apply a PCA transformation of `good_data` using `pca.transform` , and assign the results to `reduced_data` .
- Apply a PCA transformation of `log_samples` using `pca.transform` , and assign the results to `pca_samples` .

```
In [13]: 1 # TODO: Apply PCA by fitting the good data with only two dimensions
2 pca = PCA(n_components=2)
3 pca.fit(good_data)
4
5 # TODO: Transform the good data using the PCA fit above
6 reduced_data = pca.transform(good_data)
7
8 # TODO: Transform Log_samples using the PCA fit above
9 pca_samples = pca.transform(log_samples)
10
11 # Create a DataFrame for the reduced data
12 reduced_data = pd.DataFrame(reduced_data, columns = ['Dimension 1', 'Dimension 2'])
```

Observation

Run the code below to see how the log-transformed sample data has changed after having a PCA transformation applied to it using only two dimensions. Observe how the values for the first two dimensions remains unchanged when compared to a PCA transformation in six dimensions.

```
In [14]: 1 # Display sample Log-data after applying PCA transformation in two dimensions
2 display(pd.DataFrame(np.round(pca_samples, 4), columns = ['Dimension 1', 'Dimension 2']))
```

	Dimension 1	Dimension 2
0	1.7793	-0.6604
1	1.8820	0.4617
2	-0.9538	-4.0931


Visualizing a Biplot

A biplot is a scatterplot where each data point is represented by its scores along the principal components. The axes are the principal components (in this case `Dimension 1` and `Dimension 2`). In addition, the biplot shows the projection of the original features along the components. A biplot can help us interpret the reduced dimensions of the data, and discover relationships between the principal components and original features.

Run the code cell below to produce a biplot of the reduced-dimension data.

```
In [15]: 1 # Create a biplot
        2 vs.biplot(good_data, reduced_data, pca)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x10aa26d8>
```



Observation

Once we have the original feature projections (in red), it is easier to interpret the relative position of each data point in the scatterplot. For instance, a point the lower right corner of the figure will likely correspond to a customer that spends a lot on 'Milk', 'Grocery' and 'Detergents_Paper', but not so much on the other product categories.

From the biplot, which of the original features are most strongly correlated with the first component? What about those that are associated with the second component? Do these observations agree with the `pca_results` plot you obtained earlier?

Answer: The observations from the biplot are exactly the same as the scatterplot. For instance, 'Detergents_Paper', 'Grocery' and 'Milk' have much bigger increment along dimension 1. 'Delicatessen' also increase positively, but with less scale. 'Fresh' and 'Frozen' produce negative increment along dimension 1. The same conclusion also follows features on dimension 2.

Clustering

In this section, you will choose to use either a K-Means clustering algorithm or a Gaussian Mixture Model clustering algorithm to identify the various customer segments hidden in the data. You will then recover specific data points from the clusters to understand their significance by transforming them back into their original dimension and scale.

Question 6

- What are the advantages to using a K-Means clustering algorithm?
- What are the advantages to using a Gaussian Mixture Model clustering algorithm?
- Given your observations about the wholesale customer data so far, which of the two algorithms will you use and why?

**** Hint: **** Think about the differences between hard clustering and soft clustering and which would be appropriate for our dataset.

Answer: K-Means is easy and fast to compute, even with large datasets.

The biggest advantage of GMM may be data points belong to clusters in a probability way, unlike it's either 'yes' or 'no' for K-Means.

K-Means is considered as a special case of GMM, I'll try GMM in this project. In addition, there are 6 dimensions, data points tend to be more difficult to distinguish from each other when dimensions are high.

Implementation: Creating Clusters

Depending on the problem, the number of clusters that you expect to be in the data may already be known. When the number of clusters is not known *a priori*, there is no guarantee that a given number of clusters best segments the data, since it is unclear what structure exists in the data — if any. However, we can quantify the "goodness" of a clustering by calculating each data point's *silhouette coefficient*. The [silhouette coefficient](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.silhouette_score.html) (http://scikit-learn.org/stable/modules/generated/sklearn.metrics.silhouette_score.html) for a data point measures how similar it is to its assigned cluster from -1 (dissimilar) to 1 (similar). Calculating the *mean* silhouette coefficient provides for a simple scoring method of a given clustering.

In the code block below, you will need to implement the following:

- Fit a clustering algorithm to the `reduced_data` and assign it to `clusterer`.
- Predict the cluster for each data point in `reduced_data` using `clusterer.predict` and assign them to `preds`.
- Find the cluster centers using the algorithm's respective attribute and assign them to `centers`.
- Predict the cluster for each sample data point in `pca_samples` and assign them `sample_preds`.
- Import `sklearn.metrics.silhouette_score` and calculate the silhouette score of `reduced_data` against `preds`.
 - Assign the silhouette score to `score` and print the result.

In [16]:

```
1 # TODO: Apply your clustering algorithm of choice to the reduced data
2 from sklearn.mixture import GaussianMixture
3 from sklearn.metrics import silhouette_score
4 for i in range(2,21):
5     clusterer = GaussianMixture(n_components=i, random_state=0).fit(reduced_data)
6
7     # TODO: Predict the cluster for each data point
8     preds = clusterer.predict(reduced_data)
9     #print(reduced_data)
10    #print(preds)
11
12    # TODO: Find the cluster centers
13    centers = clusterer.means_
14    #print(centers)
15
16    # TODO: Predict the cluster for each transformed sample data point
17    sample_preds = clusterer.predict(pca_samples)
18    #print(sample_preds)
19
20    # TODO: Calculate the mean silhouette coefficient for the number of clusters chosen
21    score = silhouette_score(reduced_data, preds)
22    print(i, score)
23
24 clusterer = GaussianMixture(n_components=3, random_state=0).fit(reduced_data)
25
26 # TODO: Predict the cluster for each data point
27 preds = clusterer.predict(reduced_data)
28 #print(reduced_data)
29 #print(preds)
30
31 # TODO: Find the cluster centers
32 centers = clusterer.means_
33 #print(centers)
34
35 # TODO: Predict the cluster for each transformed sample data point
36 sample_preds = clusterer.predict(pca_samples)
37 #print(sample_preds)
38
39 # TODO: Calculate the mean silhouette coefficient for the number of clusters chosen
40 score = silhouette_score(reduced_data, preds)
```

```
(2, 0.446753526944537)
(3, 0.35256124857585935)
(4, 0.31513757092009265)
(5, 0.31356874310508004)
(6, 0.338434872317811)
(7, 0.24986621117251637)
(8, 0.3199219347757431)
(9, 0.3437726426497719)
(10, 0.31683966634168503)
(11, 0.3309078764672207)
(12, 0.3039499923708055)
(13, 0.3172756635940487)
(14, 0.3130072884018276)
(15, 0.3286182870215958)
(16, 0.3106885037633965)
(17, 0.29959345027936557)
(18, 0.3086582527824385)
(19, 0.2945057551462254)
(20, 0.2966919043618593)
```

Question 7

- Report the silhouette score for several cluster numbers you tried.
- Of these, which number of clusters has the best silhouette score?


Answer: I've printed the silhouette scores for cluster numbers from 2 to 20. See the above output.

When the number of clusters is 2, 3, 6, or 9, the score is slightly higher. Best core is preduced when the cluster number is 2, but in my opinion, there isn't much insight if we just divide the datasets into 2 clusters. 6 or 9 clusters may be too many. So 3 could be a balance point.

Cluster Visualization

Once you've chosen the optimal number of clusters for your clustering algorithm using the scoring metric above, you can now visualize the results by executing the code block below. Note that, for experimentation purposes, you are welcome to adjust the number of clusters for your clustering algorithm to see various visualizations. The final visualization provided should, however, correspond with the optimal number of clusters.

```
In [17]: 1 # Display the results of the clustering from implementation
        2 vs.cluster_results(reduced_data, preds, centers, pca_samples)
```



Implementation: Data Recovery

Each cluster present in the visualization above has a central point. These centers (or means) are not specifically data points from the data, but rather the *averages* of all the data points predicted in the respective clusters. For the problem of creating customer segments, a cluster's center point corresponds to *the average customer of that segment*. Since the data is currently reduced in dimension and scaled by a logarithm, we can recover the representative customer spending from these data points by applying the inverse transformations.

In the code block below, you will need to implement the following:


- Apply the inverse transform to `centers` using `pca.inverse_transform` and assign the new centers to `log_centers`.
- Apply the inverse function of `np.log` to `log_centers` using `np.exp` and assign the true centers to `true_centers`.

```
In [18]: 1 # TODO: Inverse transform the centers
        2 #print(centers)
        3 log_centers = pca.inverse_transform(centers)
        4 #print(log_centers)
        5
        6 # TODO: Exponentiate the centers
        7 true_centers = np.power(np.e, log_centers)
        8
        9 # Display the true centers
       10 segments = ['Segment {}'.format(i) for i in range(0, len(centers))]
       11 true_centers = pd.DataFrame(np.round(true_centers), columns = data.keys())
       12 true_centers.index = segments
       13 display(true_centers)
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
Segment 0	15302.0	2697.0	3323.0	3598.0	437.0	1258.0
Segment 1	5183.0	7936.0	11834.0	1069.0	4696.0	1115.0
Segment 2	6529.0	1709.0	2224.0	1496.0	292.0	567.0


```
In [19]: 1 # Add the median values and plot the 3 segments with median
2 true_centers_mean = true_centers.append(data.describe().loc['50%'])
3 true_centers_mean.plot(kind = 'bar', figsize = (16, 4))
```

Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0xefeef28>



Question 8

- Consider the total purchase cost of each product category for the representative data points above, and reference the statistical description of the dataset at the beginning of this project (specifically looking at the mean values for the various feature points). What set of establishments could each of the customer segments represent?

Hint: A customer who is assigned to 'Cluster X' should best identify with the establishments represented by the feature set of 'Segment X'. Think about what each segment represents in terms their values for the feature points chosen. Reference these values with the mean values to get some perspective into what kind of establishment they represent.

Answer: The plot from previous section shows the centroid of the 3 clusters and global median.

Segment 0 consumes 50% more 'Fresh' than global median, but less 'Detergents_Paper', maybe it's a market. Segment 1 purchases double the 'Milk', 'Grocery', and 'Detergents_Paper' of global median. It could be a retailer or hotel. Segment 2 takes all products less than median, my best guess it's a small restaurant.

Question 9

- For each sample point, which customer segment from **Question 8** best represents it?
- Are the predictions for each sample point consistent with this?

Run the code block below to find which cluster each sample point is predicted to be.

```
In [20]: 1 # Display the predictions
2 for i, pred in enumerate(sample_preds):
3     print "Sample point", i, "predicted to be in Cluster", pred
```

```
Sample point 0 predicted to be in Cluster 1
Sample point 1 predicted to be in Cluster 1
Sample point 2 predicted to be in Cluster 2
```

Answer: They are best represented by segment 1, 1, 2, respectively.

Yes, based on the earlier plot, the predication is consistent.

Conclusion

In this final section, you will investigate ways that you can make use of the clustered data. First, you will consider how the different groups of customers, the **customer segments**, may be affected differently by a specific delivery scheme. Next, you will consider how giving a label to each customer (which *segment* that customer belongs to) can provide for additional features about the customer data. Finally, you will compare the **customer segments** to a hidden variable present in the data, to see whether the clustering identified certain relationships.

Question 10

Companies will often run [A/B tests](https://en.wikipedia.org/wiki/A/B_testing) (https://en.wikipedia.org/wiki/A/B_testing) when making small changes to their products or services to determine whether making that change will affect its customers positively or negatively. The wholesale distributor is considering changing its delivery service from currently 5 days a week to 3 days a week. However, the distributor will only make this change in delivery service for customers that react positively.

- How can the wholesale distributor use the customer segments to determine which customers, if any, would react positively to the change in delivery service?*

Hint: Can we assume the change affects all customers equally? How can we determine which group of customers it affects the most?

Answer: The effect of change varies from customers. Some customers are expecting fresh food every day and may not have the capability to buffer the stock for a long time.

The A/B test can be started from one customer segment or performing concurrently for each segment. For example, we can choose 20% of the customers from each segment and treat them as the version B, meaning changing the delivery service from 5 days a week to 3 days a week for them. The rest 80% remain the same, version A. The metrics of version A and B can be obtained by observing the rate of complaints, the requests of adhoc delivery or customer survey. In this way, we will be able to know whether the 3 days a week is more preferred by different segments of customers. Decisions can be made for each customer segment instead of the entire customer base.

Question 11

Additional structure is derived from originally unlabeled data when using clustering techniques. Since each customer has a **customer segment** it best identifies with (depending on the clustering algorithm applied), we can consider '*customer segment*' as an **engineered feature** for the data. Assume the wholesale distributor recently acquired ten new customers and each provided estimates for anticipated annual spending of each product category. Knowing these estimates, the wholesale distributor wants to classify each new customer to a **customer segment** to determine the most appropriate delivery service.

- How can the wholesale distributor label the new customers using only their estimated product spending and the **customer segment** data?

Hint: A supervised learner could be used to train on the original customers. What would be the target variable?


Answer: Use the original customer data as the input and engineered feature as the output to train a supervised learning model. Then fit in the new customer spending data to predict its segment. The delivery service will be determined by the segment of the new customers.

Visualizing Underlying Distributions

At the beginning of this project, it was discussed that the 'Channel' and 'Region' features would be excluded from the dataset so that the customer product categories were emphasized in the analysis. By reintroducing the 'Channel' feature to the dataset, an interesting structure emerges when considering the same PCA dimensionality reduction applied earlier to the original dataset.

Run the code block below to see how each data point is labeled either 'HoReCa' (Hotel/Restaurant/Cafe) or 'Retail' the reduced space. In addition, you will find the sample points are circled in the plot, which will identify their labeling.

```
In [21]: 1 # Display the clustering results based on 'Channel' data
        2 vs.channel_results(reduced_data, outliers, pca_samples)
```



Question 12

- How well does the clustering algorithm and number of clusters you've chosen compare to this underlying distribution of Hotel/Restaurant/Cafe customers to Retailer customers?
- Are there customer segments that would be classified as purely 'Retailers' or 'Hotels/Restaurants/Cafes' by this distribution?
- Would you consider these classifications as consistent with your previous definition of the customer segments?

Answer: It seems I've made the things over complicated that I chose 9 clusters.

No, there are some green dots in red region, also some red dots in green region.

No, because the 2 clusters are not perfectly separated, the final converged result depends on the initial cluster center selection.

Note: Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.